

GENERATIVE AI GOVERNANCE

Real-World Failures & Lessons, Explained Simply | 2023 – 2026

Every case is a true story with a clear source you can look up.
Every slide ends with one plain-language lesson.

GenAI for Software Professionals curriculum

Updated June 2026 |



WHAT WE'LL LOOK AT — SIX KINDS OF THINGS THAT GO WRONG

01

Privacy

Using people's personal data without permission

02

Unfair Treatment (Bias)

AI that discriminates against certain people

03

Fake Content (Deepfakes)

AI-made images and videos that fool people

04

Security

Automated systems causing big breakdowns

05

Copyright

AI trained on others' work without paying

06

Lessons & Rules

What to do about it + the new laws

ChatGPT's maker was fined for mishandling people's data

WHAT HAPPENED (in plain words)

- Italy's privacy watchdog fined OpenAI €15 million (about ₹140 crore) in Dec 2024.
- Why? OpenAI used people's personal information to train ChatGPT without a proper legal reason.
- It also didn't tell the regulator about a 2023 data breach, and had no way to check if users were children.
- Europe's privacy law (called GDPR) requires companies to be open about how they use your data.

THE LESSON

Before you use people's personal data to build or train an AI, get permission and a clear legal reason — and be honest about what you collect.

Plain takeaway:

"If you wouldn't be comfortable explaining how you got the data, don't use it."

2026 update: An Italian court cancelled the fine in March 2026 — proof that the rules around AI are still being worked out in courts.

Source: *The Hacker News (Dec 2024) & Cross-Border Data Forum (Mar 2026) — full link on References slide [1][19]*

A hiring AI was accused of rejecting people unfairly

The story: A man named Derek Mobley applied for hundreds of jobs through screening software made by a company called Workday. He was turned down almost every time. He sued, saying the AI was rejecting him because of his age, race, and a disability — not his skills.

WHY THIS ONE IS A BIG DEAL

- In 2025 a US court allowed it to become a group lawsuit — so potentially millions of rejected applicants could join.
- It's the first big case where the AI tool's maker (not just the employer) could be held responsible.
- A 2024 University of Washington study found AI screening tools favoured white-sounding names about 85% of the time.

The lesson: Test your AI for fairness before and after you launch it, and keep a human in the loop for big decisions like hiring.

AI-made fakes spread during elections around the world

80+

countries saw AI fakes in their 2024 elections

What's a deepfake? A photo, video, or voice clip made by AI that looks real but isn't. It can show someone saying or doing something they never did — which is why it's so easy to fool people during an election.

United States (2024)

Fake AI images showed pop star Taylor Swift appearing to endorse a candidate. She had not — and later spoke out about it.

Source: CIGI / R Street Institute, 2024–25 — References [3]

India (2024)

Fake videos showed famous people seeming to criticise or back political parties; they spread fast on WhatsApp.

Source: Brennan Center for Justice, 2025 — References [3]

The lesson: Label AI-made content, and double-check surprising images or videos before believing or sharing them.

One automatic update knocked out computers worldwide

\$5–10 Billion

estimated cost to businesses worldwide

What happened (plain): In 2024 a security company (CrowdStrike) pushed one automatic software update to computers everywhere at once. It had a fault — and millions of machines crashed, grounding flights and stopping hospitals and banks.

WHY IT'S A GOVERNANCE STORY

- The update went to everyone at once — there was no "try it on a few machines first" safety step.
- There was no quick "off switch" to stop it once things started going wrong.
- As more tasks get handed to automated systems, one bad change can cause damage on a huge scale.

The lesson: Roll out automated changes gradually, and always keep a way to stop them fast if something breaks.

Was it fair to train AI on other people's work?

AI tools learn by studying huge amounts of text, images, and music. Often that material is owned by someone. By 2025, more than 70 lawsuits had been filed. A few key ones, in plain terms:

The New York Times vs. OpenAI

The newspaper says ChatGPT was trained on its articles without permission. A landmark test of what's allowed.

Source: *Copyright Alliance, 2025 — Ref [5]*

Music labels vs. Suno & Udio

Record companies say these AI music apps copied their songs to learn how to make new ones.

Source: *McKool Smith AI tracker, 2024 — Ref [20]*

Thomson Reuters vs. Ross

A court decided copying a legal database to train an AI was NOT fair — an early win for the owner.

Source: *Norton Rose Fulbright, 2025 — Ref [13]*

Disney, Universal & Warner vs. Midjourney

Big Hollywood studios say this AI makes images of their characters (Shrek, Star Wars) without a licence.

Source: *The Hollywood Reporter, 2025 — Ref [18]*

The lesson: Keep records of what your AI was trained on, get licences where you can, and avoid pirated material.

Three times a company's own chatbot caused trouble

Air Canada

2024

What went wrong: Its chatbot made up a refund rule that didn't exist. A customer relied on it and was refused.

Result: A tribunal ruled the airline had to honour it — you can't blame your own chatbot.

Source: *American Bar Association (Moffatt v. Air Canada) — Ref [7]*

New York City

2024

What went wrong: The city's official business-help chatbot gave advice that would have had owners breaking the law.

Result: Public embarrassment and lost trust; it had gone live without proper checks.

Source: *Awesm AI governance review, 2026 — Ref [6]*

DPD (delivery)

2024

What went wrong: After a software update, the delivery chatbot swore at a customer and wrote a poem mocking its own company.

Result: It went viral (1M+ views); DPD switched the AI part off to fix it.

Source: *BBC / TIME, Jan 2024 — Ref [12]*

The lesson: Your chatbot speaks for you. Test it, set limits, and check it before customers ever see it.

BONUS — How the worries began (2023)

Staff pasting company secrets into chatbots

The wake-up call (Samsung, 2023): Engineers pasted secret company code and meeting notes into ChatGPT to get help. That information left the company and couldn't be pulled back. Samsung banned the tool soon after.

OTHER COMPANIES THAT REACTED THE SAME WAY

Apple

Limited staff use of ChatGPT and Copilot

JPMorgan

Restricted use over compliance worries

Amazon

Told staff: don't paste in code or secrets

Goldman & Citi

Among first banks to block it

Verizon

Blocked it on company systems

Many more

Moved to "ask-first" approval rules

Eye-opener: In a 2023 survey, 68% of staff using ChatGPT at work were doing so without telling their employer. This is called "Shadow AI" — AI use nobody is keeping track of.

Source: CNBC (May 2023) & Fortune (Apple) — References [9]. The lesson: have a clear rule on what staff may share with AI tools.

THE BIG LESSONS — in one line each

Experts looked at hundreds of 2025 incidents. Most problems weren't about clever technology — they were about people, ownership, and rules.

AI is sometimes confidently wrong

It can state false things as if they're true ("hallucinations"). Always have a person check important answers.

Know who owns each AI system

Most failures trace back to a simple gap: nobody clearly owned it or watched over it. Name an owner.

Treat AI tools like trusted staff

Give them only the access they need, keep logs, and review them — same as you would a new employee.

You can't manage what you can't see

Keep a simple list of every AI tool in use. "Shadow AI" (hidden use) is the biggest blind spot.

Start with the goal, not the gadget

Decide what result you want first. Don't adopt AI just because it's new and exciting.

Good governance is an advantage

Companies that stay organised and accountable build more trust — and win more business.

Europe's AI law starts being enforced

AUG 2

2026

Main rules for higher-risk AI take full effect

Real penalties — Fines up to €35 million, or 7% of a company's worldwide sales — whichever is larger.

Be upfront about AI — Companies must tell people when they're talking to AI, and clearly label AI-made content.

Mark AI-made content — AI-made text, images, audio, and video must carry a hidden "made by AI" marker (from Dec 2026).

Some AI is simply banned — Uses seen as too risky — like government "social scoring" of citizens — are not allowed at all.

IT'S NOT JUST EUROPE

26 major AI companies (including Microsoft, Google, Amazon, OpenAI and Anthropic) signed a shared code of good practice. Other countries are following with their own rules. The hard part that's still unsolved: enforcing rules when AI crosses borders.

YOUR SIMPLE 6-STEP CHECKLIST

How to use AI responsibly — drawn straight from the cases above

1

Put someone in charge

Name a small team (legal, tech, risk) that owns AI decisions.

2

List every AI tool you use

You can't manage what you don't know about. Keep it written down.

3

Protect people's data

Get permission, have a legal reason, and report breaches fast.

4

Check for fairness

Test that the AI treats everyone equally, especially in big decisions.

5

Respect copyright

Track what your AI learned from; get licences; avoid pirated data.

6

Add safety brakes

Roll out changes slowly and keep an "off switch" ready.

REMEMBER THESE SIX THINGS



Handle data with care

OpenAI was fined €15M for using data the wrong way. Permission first.



Fair AI is the law

The Workday case shows hiring AI can land you in a group lawsuit.



Fakes are everywhere

AI deepfakes hit elections in 80+ countries. Label and verify content.



Build in safety brakes

One bad auto-update cost businesses billions. Roll out slowly.



Respect others' work

70+ copyright lawsuits. Track your training data; get licences.



Stay organised

Most failures came from no clear owner. Good governance wins trust.

SOURCES — LOOK UP ANY CASE YOURSELF

[1] Privacy — OpenAI fine (The Hacker News, 2024) <https://thehackernews.com/2024/12/italy-fines-openai-15-million-for.html>

[2] Bias — Mobley v. Workday (Lathrop GPM) <https://www.lathropgpm.com/insights/lawsuits-alleging-systemic-bias-in-ai-algorithmic-screening-tools>

[3] Deepfakes — AI in elections (CIGI) <https://www.cigionline.org/articles/then-and-now-how-does-ai-electoral-interference-compare-in-2025/>

[4] Security — CrowdStrike / machine identity (arXiv) <https://arxiv.org/pdf/2604.06148>

[5] Copyright overview 2025 (Copyright Alliance) <https://copyrightalliance.org/ai-copyright-lawsuit-developments-2025/>

[6] NYC chatbot — governance failures (Awesml) <https://awesml.com/tech/ai-governance-failures>

[7] Air Canada chatbot ruling (American Bar Assn.) https://www.americanbar.org/groups/business_law/resources/business-law-today/2024-february/bc-tribunal-confirms-companies-remain-liable-information-provided-ai-chatbot/

[8] AI as a governance issue (Monash Lens) <https://lens.monash.edu/responsible-ai-is-now-a-governance-risk-not-an-ethics-debate>

[9] Data leaks — Samsung bans ChatGPT (CNBC) <https://www.cnn.com/2023/05/02/samsung-bans-use-of-ai-like-chatgpt-for-staff-after-misuse-of-chatbot.html>

[10] Lessons — Avoiding AI Pitfalls in 2026 (ISACA) <https://www.isaca.org/resources/news-and-trends/isaca-now-blog/2025/avoiding-ai-pitfalls-in-2026>

[11] Governance lessons 2025 (Nat'l Law Review) <https://natlawreview.com/article/ai-governance-programs-provide-competitive-advantage>

[12] DPD chatbot incident (TIME) <https://time.com/6564726/ai-chatbot-dpd-curses-criticizes-company/>

[13] Copyright 2026 update (Norton Rose Fulbright) <https://www.nortonrosefulbright.com/en/knowledge/publications/ce8eaa5f>

[14] Hiring bias resume study (UW / Quinn Emanuel) <https://www.quinnemanuel.com/the-firm/publications/when-machines-discriminate-the-rise-of-ai-bias-lawsuits/>

[15] Singapore Model AI Governance Framework <https://www.pdpc.gov.sg/Help-and-Resources/2020/01/Model-AI-Governance-Framework>

[16] EU AI Act — official (European Commission) <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>

[17] EU AI Act 2026 deadlines & fines (Latham) <https://www.lw.com/en/insights/ai-act-update-eu-resolves-to-change-rules-and-extend-deadlines>

[18] Disney/Universal v. Midjourney (Hollywood Reporter) <https://www.hollywoodreporter.com/business/business-news/disney-universal-midjourney-1236262563/>

[19] OpenAI fine cancelled 2026 (Cross-Border Data) <https://www.crossborderdataforum.org/generative-ai-and-gdpr-enforcement-in-europe-a-lot-of-noise-one-fine-zero-survivors/>

[20] Music labels v. Suno & Udio (McKool Smith) <https://www.mckoolsmith.com/newsroom-ai-litigation-36>